

K-Nearest Neighbors
Naive Bayes
Decision Tree
Random Forest
XgBoost
Support Vector Machine
Neural Network

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
```

```
import os
```

```
import warnings
warnings.filterwarnings('ignore')
# read file
voice=pd.read_csv('/voice.csv')
voice.head()
```

	meanfreq	sd	median	Q25	Q75	IQR	skew	kurt	sp.ent	sfm	...	centroid	meanfun	minfun
0	0.059781	0.064241	0.032027	0.015071	0.090193	0.075122	12.863462	274.402906	0.893369	0.491918	...	0.059781	0.084279	0.015702
1	0.066009	0.067310	0.040229	0.019414	0.092666	0.073252	22.423285	634.613855	0.892193	0.513724	...	0.066009	0.107937	0.015826
2	0.077316	0.083829	0.036718	0.008701	0.131908	0.123207	30.757155	1024.927705	0.846389	0.478905	...	0.077316	0.098706	0.015656
3	0.151228	0.072111	0.158011	0.096582	0.207955	0.111374	1.232831	4.177296	0.963322	0.727232	...	0.151228	0.088965	0.017798
4	0.135120	0.079146	0.124656	0.078720	0.206045	0.127325	1.101174	4.333713	0.971955	0.783568	...	0.135120	0.106398	0.016931

5 rows × 21 columns



```
voice.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3168 entries, 0 to 3167
Data columns (total 21 columns):
#   Column      Non-Null Count  Dtype
---  -
0   meanfreq    3168 non-null   float64
1   sd          3168 non-null   float64
2   median      3168 non-null   float64
3   Q25         3168 non-null   float64
4   Q75         3168 non-null   float64
5   IQR         3168 non-null   float64
6   skew        3168 non-null   float64
7   kurt        3168 non-null   float64
8   sp.ent      3168 non-null   float64
9   sfm         3168 non-null   float64
10  mode        3168 non-null   float64
11  centroid    3168 non-null   float64
12  meanfun     3168 non-null   float64
13  minfun      3168 non-null   float64
14  maxfun      3168 non-null   float64
15  meandom     3168 non-null   float64
16  mindom      3168 non-null   float64
17  maxdom      3168 non-null   float64
18  dfrange     3168 non-null   float64
19  modindx     3168 non-null   float64
20  label       3168 non-null   object
dtypes: float64(20), object(1)
memory usage: 519.9+ KB
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

voice.describe()

	meanfreq	sd	median	Q25	Q75	IQR	skew	kurt	sp.ent	sfm	
count	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000
mean	0.180907	0.057126	0.185621	0.140456	0.224765	0.084309	3.140168	36.568461	0.895127	0.408216	
std	0.029918	0.016652	0.036360	0.048680	0.023639	0.042783	4.240529	134.928661	0.044980	0.177521	
min	0.039363	0.018363	0.010975	0.000229	0.042946	0.014558	0.141735	2.068455	0.738651	0.036876	
25%	0.163662	0.041954	0.169593	0.111087	0.208747	0.042560	1.649569	5.669547	0.861811	0.258041	
50%	0.184838	0.059155	0.190032	0.140286	0.225684	0.094280	2.197101	8.318463	0.901767	0.396335	
75%	0.199146	0.067020	0.210618	0.175939	0.243660	0.114175	2.931694	13.648905	0.928713	0.533676	
max	0.251124	0.115273	0.261224	0.247347	0.273469	0.252225	34.725453	1309.612887	0.981997	0.842936	



Preprocessing: label encoder and normalization

```
from sklearn import preprocessing
le = preprocessing.LabelEncoder()
voice["label"] = le.fit_transform(voice["label"])
le.classes_

array(['female', 'male'], dtype=object)

voice[:]=preprocessing.MinMaxScaler().fit_transform(voice)
voice.head()
```

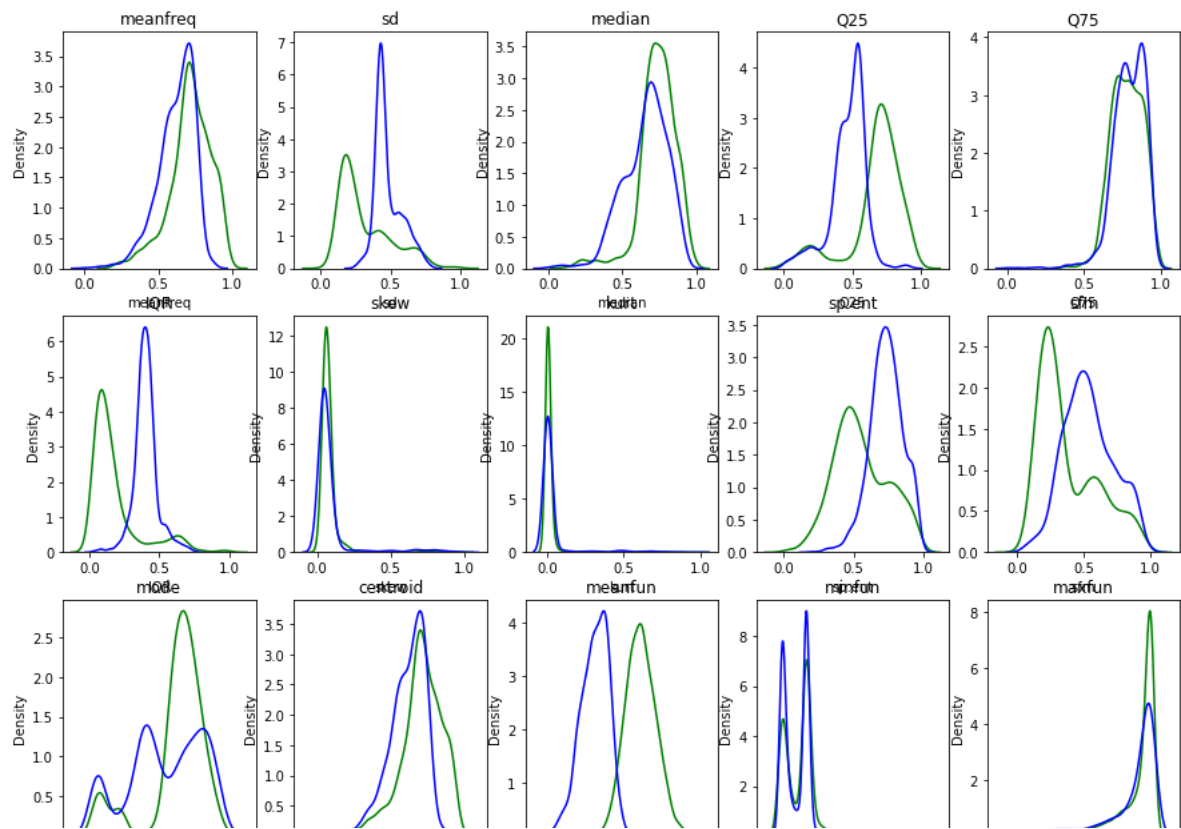
	meanfreq	sd	median	Q25	Q75	IQR	skew	kurt	sp.ent	sfm	...	centroid	meanfun	minfun	m
0	0.096419	0.473409	0.084125	0.060063	0.204956	0.254828	0.367853	0.208279	0.635798	0.564526	...	0.096419	0.157706	0.030501	0.9:
1	0.125828	0.505075	0.116900	0.077635	0.215683	0.246961	0.644279	0.483766	0.630964	0.591578	...	0.125828	0.287642	0.031140	0.8:
2	0.179222	0.675536	0.102873	0.034284	0.385912	0.457148	0.885255	0.782275	0.442738	0.548382	...	0.179222	0.236945	0.030264	0.9:
3	0.528261	0.554611	0.587559	0.389906	0.715802	0.407358	0.031549	0.001613	0.923261	0.856457	...	0.528261	0.183442	0.041287	0.8:
4	0.452195	0.627209	0.454272	0.317627	0.707515	0.474474	0.027742	0.001732	0.958736	0.926348	...	0.452195	0.279190	0.036829	0.9:

5 rows x 21 columns



Visualization

```
import seaborn as sns
import matplotlib.pyplot as plt
plt.subplots(4,5,figsize=(15,15))
for i in range(1,21):
    plt.subplot(4,5,i)
    plt.title(voice.columns[i-1])
    sns.kdeplot(voice.loc[voice['label'] == 0, voice.columns[i-1]], color= 'green', label='F')
    sns.kdeplot(voice.loc[voice['label'] == 1, voice.columns[i-1]], color= 'blue', label='M')
```



At first glance, most significant features are Q25, IQR and meanfun. We will build models by using the 20 features and the 3 distinct features.

Using K-Nearest Neighbors, Naive Bayes, Decision Tree, Random Forest, XgBoost, Support Vector Machine, Neural Network to build models

```
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

from sklearn import neighbors
from sklearn import naive_bayes
from sklearn import tree
from sklearn import ensemble
from sklearn import svm
from sklearn import neural_network
import xgboost

# Split the data
train, test = train_test_split(voice, test_size=0.3)

train.head()
```

	meanfreq	sd	median	Q25	Q75	IQR	skew	kurt	sp.ent	sfm	...	centroid	meanfun	minfun
553	0.632757	0.524496	0.653215	0.442806	0.870848	0.502739	0.061808	0.005041	0.773748	0.650084	...	0.632757	0.198393	0.050357
216	0.652996	0.425988	0.654083	0.504866	0.803398	0.372790	0.030166	0.001605	0.826119	0.585297	...	0.652996	0.354916	0.047256
2373	0.945598	0.139861	0.936918	0.933495	0.934897	0.054662	0.056913	0.003513	0.299454	0.130639	...	0.945598	0.610983	0.197717
316	0.558356	0.445362	0.577121	0.419390	0.711356	0.372389	0.056048	0.004304	0.800397	0.636284	...	0.558356	0.296397	0.022796
3058	0.801846	0.269198	0.819545	0.795136	0.819254	0.086357	0.062483	0.004624	0.517888	0.283671	...	0.801846	0.821097	0.088787

5 rows × 15 columns

```
x_train = train.iloc[:, :-1]
y_train = train["label"]
```

```

x_test = test.iloc[:, :-1]
y_test = test["label"]

x_train3 = train[["meanfun", "IQR", "Q25"]]
y_train3 = train["label"]
x_test3 = test[["meanfun", "IQR", "Q25"]]
y_test3 = test["label"]

def classify(model, x_train, y_train, x_test, y_test):
    from sklearn.metrics import classification_report
    target_names = ['female', 'male']
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    print(classification_report(y_test, y_pred, target_names=target_names, digits=4))

```

▼ K-Nearest Neighbors

Using `neighbors.KNeighborsClassifier()` to build the model.

```


def knn_error(k, x_train, y_train, x_test, y_test):
    error_rate = []
    K = range(1, k)
    for i in K:
        knn = neighbors.KNeighborsClassifier(n_neighbors = i)
        knn.fit(x_train, y_train)
        y_pred = knn.predict(x_test)
        error_rate.append(np.mean(y_pred != y_test))
    kloc = error_rate.index(min(error_rate))
    print("Lowest error is %s occurs at k=%s." % (error_rate[kloc], K[kloc]))

plt.plot(K, error_rate, color='blue', linestyle='dashed', marker='o',
         markerfacecolor='red', markersize=10)
plt.title('Error Rate vs. K Value')
plt.xlabel('K')
plt.ylabel('Error Rate')
plt.show()
return K[kloc]

```

```
k=knn_error(21, x_train, y_train, x_test, y_test)
```

Lowest error is 0.026288117770767613 occurs at k=4.
Error Rate vs. K Value



```


model = neighbors.KNeighborsClassifier(n_neighbors = k)
classify(model, x_train, y_train, x_test, y_test)

```

	precision	recall	f1-score	support
female	0.9708	0.9769	0.9739	477
male	0.9766	0.9705	0.9735	474
accuracy			0.9737	951
macro avg	0.9737	0.9737	0.9737	951
weighted avg	0.9737	0.9737	0.9737	951

```
k=knn_error(21, x_train3, y_train3, x_test3, y_test3)
```

Lowest error is 0.027339642481598318 occurs at k=3.



```
model = neighbors.KNeighborsClassifier(n_neighbors = k)
classify(model,x_train,y_train,x_test,y_test)
```

	precision	recall	f1-score	support
female	0.9787	0.9644	0.9715	477
male	0.9647	0.9789	0.9717	474
accuracy			0.9716	951
macro avg	0.9717	0.9716	0.9716	951
weighted avg	0.9717	0.9716	0.9716	951

Naive Bayes

Using `naive_bayes.GaussianNB()` to build the model.

```
model=naive_bayes.GaussianNB()
classify(model,x_train,y_train,x_test,y_test)
```

	precision	recall	f1-score	support
female	0.8645	0.8826	0.8734	477
male	0.8793	0.8608	0.8699	474
accuracy			0.8717	951
macro avg	0.8719	0.8717	0.8717	951
weighted avg	0.8719	0.8717	0.8717	951

```
model=naive_bayes.GaussianNB()
classify(model,x_train3,y_train3,x_test3,y_test3)
```

	precision	recall	f1-score	support
female	0.9661	0.9560	0.9610	477
male	0.9562	0.9662	0.9612	474
accuracy			0.9611	951
macro avg	0.9611	0.9611	0.9611	951
weighted avg	0.9611	0.9611	0.9611	951

Decision Tree

Using `tree.DecisionTreeClassifier()` to build the model.

```
#Find the best parameter to prune the tree
def dt_error(n,x_train,y_train,x_test,y_test):
    nodes = range(2, n)
    error_rate = []
    for k in nodes:
        model = tree.DecisionTreeClassifier(max_leaf_nodes=k)
        model.fit(x_train, y_train)
        y_pred = model.predict(x_test)
        error_rate.append(np.mean(y_pred != y_test))
    kloc = error_rate.index(min(error_rate))
    print("Lowest error is %s occurs at n=%s." % (error_rate[kloc], nodes[kloc]))
    plt.plot(nodes, error_rate, color='blue', linestyle='dashed', marker='o',
```


```

        markerfacecolor='red', markersize=10)
plt.xlabel('Tree Size')
plt.ylabel('Cross-Validated MSE')
plt.show()
return nodes[kloc]

```

```
n=dt_error(10,x_train,y_train,x_test,y_test)
```

Lowest error is 0.033648790746582544 occurs at n=7.



```


#prune tree
pruned_tree = tree.DecisionTreeClassifier(criterion = 'gini', max_leaf_nodes = n)
classify(pruned_tree,x_train,y_train,x_test,y_test)

```

	precision	recall	f1-score	support
female	0.9704	0.9623	0.9663	477
male	0.9623	0.9705	0.9664	474
accuracy			0.9664	951
macro avg	0.9664	0.9664	0.9664	951
weighted avg	0.9664	0.9664	0.9664	951

```
n=dt_error(15,x_train3,y_train3,x_test3,y_test3)
```

Lowest error is 0.03259726603575184 occurs at n=7.



```

#prune tree
pruned_tree = tree.DecisionTreeClassifier(criterion = 'gini', max_leaf_nodes = n)
classify(pruned_tree,x_train3,y_train3,x_test3,y_test3)

```

	precision	recall	f1-score	support
female	0.9705	0.9644	0.9674	477
male	0.9644	0.9705	0.9674	474
accuracy			0.9674	951
macro avg	0.9674	0.9674	0.9674	951
weighted avg	0.9674	0.9674	0.9674	951

▼ Random Forest


Using `ensemble.RandomForestClassifier()` to build the model.

```
def rf_error(n,x_train,y_train,x_test,y_test):
    error_rate = []
    e=range(1,n,20)
    for i in e:
        model = ensemble.RandomForestClassifier(n_estimators = i)
        model.fit(x_train, y_train)
        y_pred = model.predict(x_test)
        error_rate.append(np.mean(y_pred != y_test))
    nloc = error_rate.index(min(error_rate))
    print("Lowest error is %s occurs at n=%s." % (error_rate[nloc], e[nloc]))

    plt.plot(e, error_rate, color='blue', linestyle='dashed', marker='o',
             markerfacecolor='red', markersize=10)
    plt.title('Error Rate vs. n Value')
    plt.xlabel('n')
    plt.ylabel('Error Rate')
    plt.show()
    return e[nloc]
```

```
e=rf_error(100,x_train,y_train,x_test,y_test)
```

Lowest error is 0.026288117770767613 occurs at n=61.
Error Rate vs. n Value




```
model=ensemble.RandomForestClassifier(n_estimators = e)
classify(model,x_train,y_train,x_test,y_test)
```

	precision	recall	f1-score	support
female	0.9748	0.9748	0.9748	477
male	0.9747	0.9747	0.9747	474
accuracy			0.9748	951
macro avg	0.9748	0.9748	0.9748	951
weighted avg	0.9748	0.9748	0.9748	951

```
e=rf_error(100,x_train3,y_train3,x_test3,y_test3)
```

Lowest error is 0.03259726603575184 occurs at n=41.



```
model=ensemble.RandomForestClassifier(n_estimators = e)
classify(model,x_train3,y_train3,x_test3,y_test3)
```

	precision	recall	f1-score	support
female	0.9705	0.9644	0.9674	477

male	0.9644	0.9705	0.9674	474
accuracy			0.9674	951
macro avg	0.9674	0.9674	0.9674	951
weighted avg	0.9674	0.9674	0.9674	951

▼ XgBoost

Using `xgboost.XGBClassifier()` to build the model.

```
model = xgboost.XGBClassifier()
classify(model,x_train,y_train,x_test,y_test)
```

	precision	recall	f1-score	support
female	0.9729	0.9790	0.9760	477
male	0.9788	0.9726	0.9757	474
accuracy			0.9758	951
macro avg	0.9758	0.9758	0.9758	951
weighted avg	0.9758	0.9758	0.9758	951

```
model = xgboost.XGBClassifier()
classify(model,x_train3,y_train3,x_test3,y_test3)
```

	precision	recall	f1-score	support
female	0.9703	0.9602	0.9652	477
male	0.9603	0.9705	0.9654	474
accuracy			0.9653	951
macro avg	0.9653	0.9653	0.9653	951
weighted avg	0.9654	0.9653	0.9653	951

▼ Support Vector Machine

Using `svm.SVC()` to build the model.

```
def svm_kernel(x_train,y_train,x_test,y_test):
    rate=[]
    kernel=['rbf','poly','linear']
    for i in kernel:
        model=svm.SVC(kernel=i).fit(x_train,y_train)
        y_pred=model.predict(x_train)
        print(i, ' in-sample accuracy in SVM: ', accuracy_score(y_train,y_pred))
        y_pred=model.predict(x_test)
        print(i, ' out-of-sample accuracy in SVM: ', accuracy_score(y_test,y_pred))
        rate.append(accuracy_score(y_test,y_pred))
    nloc = rate.index(max(rate))
    print("Highest accuracy is %s occurs at %s kernel." % (rate[nloc], kernel[nloc]))
    return kernel[nloc]
```



```
def svm_error(k,C,x_train,y_train,x_test,y_test):
    error_rate = []
    C=range(1,C)
    for i in C:
        model=svm.SVC(kernel=k,C=i).fit(x_train,y_train)
        model.fit(x_train, y_train)
        y_pred = model.predict(x_test)
        error_rate.append(np.mean(y_pred != y_test))
    cloc = error_rate.index(min(error_rate))
    print("Lowest error is %s occurs at C=%s." % (error_rate[cloc], C[cloc]))


    plt.plot(C, error_rate, color='blue', linestyle='dashed', marker='o',
             markerfacecolor='red', markersize=10)
    plt.title('Error Rate vs. C Value')
    plt.xlabel('C')
    plt.ylabel('Error Rate')
    plt.show()
    return C[cloc]
```

```
k=svm_kernel(x_train,y_train,x_test,y_test)
```

```
rbf in-sample accuracy in SVM: 0.979702300405954
rbf out-of-sample accuracy in SVM: 0.9674027339642481
poly in-sample accuracy in SVM: 0.9828597203428056
poly out-of-sample accuracy in SVM: 0.9747634069400631
linear in-sample accuracy in SVM: 0.9769959404600812
linear out-of-sample accuracy in SVM: 0.9726603575184016
Highest accuracy is 0.9747634069400631 occurs at poly kernel.
```

```
c=svm_error(k,10,x_train,y_train,x_test,y_test)
```

Lowest error is 0.025236593059936908 occurs at C=1.



```
model=svm.SVC(kernel=k,C=c)
```

```
classify(model,x_train,y_train,x_test,y_test)
```


	precision	recall	f1-score	support
female	0.9729	0.9769	0.9749	477
male	0.9767	0.9726	0.9746	474
accuracy			0.9748	951
macro avg	0.9748	0.9748	0.9748	951
weighted avg	0.9748	0.9748	0.9748	951

```
k=svm_kernel(x_train3,y_train3,x_test3,y_test3)
```

```
rbf in-sample accuracy in SVM: 0.9760938204781235
rbf out-of-sample accuracy in SVM: 0.9621451104100947
poly in-sample accuracy in SVM: 0.9738385205232296
poly out-of-sample accuracy in SVM: 0.9631966351209253
linear in-sample accuracy in SVM: 0.9693279206134416
linear out-of-sample accuracy in SVM: 0.9589905362776026
Highest accuracy is 0.9631966351209253 occurs at poly kernel.
```

```
c=svm_error(k,10,x_train3,y_train3,x_test3,y_test3)
```

Lowest error is 0.03680336487907466 occurs at C=1.



```
model=svm.SVC(kernel=k,C=c)
classify(model,x_train3,y_train3,x_test3,y_test3)
```

	precision	recall	f1-score	support
female	0.9722	0.9539	0.9630	477
male	0.9545	0.9726	0.9634	474
accuracy			0.9632	951
macro avg	0.9633	0.9632	0.9632	951
weighted avg	0.9634	0.9632	0.9632	951

Neural Network

Using neural_network.MLPClassifier to build the model.


```
def nn_error(n,x_train,y_train,x_test,y_test):
    error_rate = []
    hidden_layer=range(1,n)
    for i in hidden_layer:
        model = neural_network.MLPClassifier(solver='adam', alpha=1e-5,
                                              hidden_layer_sizes=i,
                                              activation='logistic',random_state=17,
                                              max_iter=2000)

        model.fit(x_train, y_train)
        y_pred = model.predict(x_test)
        error_rate.append(np.mean(y_pred != y_test))
    kloc = error_rate.index(min(error_rate))
    print("Lowest error is %s occurs at C=%s." % (error_rate[kloc], hidden_layer[kloc]))

    plt.plot(hidden_layer, error_rate, color='blue', linestyle='dashed', marker='o',
             markerfacecolor='red', markersize=10)
    plt.title('Error Rate vs. Hidden Layer Size')
    plt.xlabel('Size')
    plt.ylabel('Error Rate')
    plt.show()
    return hidden_layer[kloc]
```

```
h=nn_error(20,x_train,y_train,x_test,y_test)
```

Lowest error is 0.030494216614090432 occurs at C=14.



```
model = neural_network.MLPClassifier(solver='adam', alpha=1e-5,
                                     hidden_layer_sizes=h,
                                     activation='logistic',random_state=17,
```

